



Per-prompt differences behind the two preregistered comparisons



One bar per prompt, each the mean absolute error over that prompt's 16 non-no-op interventions. The dashed line is the overall point estimate; the interval in each title is the preregistered paired bootstrap over 32 prompt groups, 10,000 resamples.